

3. Exercise

Timo Bergerbusch 344408

Thomas N  veke 311392

Shu Xing 381176

05.11.2017

Exercise 3.1

1. a)

$$\pi_{A,B}(R) \bowtie \pi_{B,C}(\pi_{A,C}(\sigma_{B=1}(R))) \bowtie \pi_{A,B}(R)$$

$$S_1 := \{[a, b] \mid \exists c R(a, b, c)\}$$

$$S_2 := \{[a, c] \mid \exists a, c R(a, 1, c)\}$$

$$\pi_{B=1}(S_2 \bowtie S_1) := \{[b, c] \mid \exists a, c_1, c_2 R(a, b, c_1) \wedge R(a, 1, c_2)\}$$

$$\text{Insgesamt} := \{[a, b, c] \mid \exists a_1, c_1, c_2 (R(a, b, c_1) \wedge R(a_1, 1, c) \wedge R(a_1, b, c_2))\}$$

1.b)

T_1			
a	b	c_1	R
a_1	1	c	R
a_1	b	c_2	R

2)

$$h_1 : T_2 \rightarrow T_1 : a \rightarrow a, b \rightarrow b, a_5 \rightarrow a_1, b_5 \rightarrow b_1, c_4 \rightarrow c_1 \\ \Rightarrow T_2 \subseteq T_1 \quad (1)$$

$$h_2 : T_1 \rightarrow T_2 : a \rightarrow a, b \rightarrow b, b_1 \rightarrow b, c_1 \rightarrow c_4, a_1 \rightarrow a, b_2 \rightarrow b_5, c_2 \rightarrow c_4, b_3 \rightarrow b_5, c_3 \rightarrow c_4 \\ \Rightarrow T_1 \subseteq T_2 \quad (2)$$

$$(1) \& (2) \Rightarrow T_1 \equiv T_2$$

Exercise 3.2

Given 16 buffer pages (B) and :

- Album has a size of 10.000 pages (M), 40 bytes record size (s_1) and 100 tuples/page (p_A)
- Track has a size of 200.000 pages (N), 30 bytes record size (s_2) and 80 tuples/page (p_T)

1)

Since the simple nested loop join is a double iteration over both relations the I/O requirements can be calculated as follows:

$$M + p_A \cdot M \cdot N = 10.000 + 100 \cdot 10.000 \cdot 200.000 = 200.000.010.000 \text{ I/Os}$$

2)

Since the block nested loop join uses 1 input and 1 output buffer the number of I/Os can be calculated with the following formula:

$$M + \lceil \frac{M}{B-2} \rceil \cdot N = 10.000 + \lceil \frac{10.000}{16-2} \rceil \cdot 200.000 = 143.060.000$$

3)

Similarities:

- double loop schema

Differences:

- usage of buffer pages
- number of *outer* elements in the memory
- hashing used in the block nested loop

Explanation:

Lets define the outer loop to be over E and the inner loop to be over T.

The *block nested loop join* loads up to $B - 2$ pages of E into the memory. These blocks get stored in a hash table. Now every input of a page from T can be used via one of the two reserved pages. After hash probing the input, if there is any match in a block of Es, we can reduce the number of scanned Es per T and therefore reduce I/Os.