

### 3. Exercise

Timo Bergerbusch 344408

Thomas N  veke 311392

Shu Xing 381176

05.11.2017

#### Exercise 3.1

##### 1. a)

$$\pi_{A,B}(R) \bowtie \pi_{B,C}(\pi_{A,C}(\sigma_{B=1}(R))) \bowtie \pi_{A,B}(R)$$

$$S_1 := \{[a, b] \mid \exists c R(a, b, c)\}$$

$$S_2 := \{[a, c] \mid \exists a, c R(a, 1, c)\}$$

$$\pi_{B=1}(S_2 \bowtie S_1) := \{[b, c] \mid \exists a, c_1, c_2 R(a, b, c_1) \wedge R(a, 1, c_2)\}$$

$$\text{Insgesamt} := \{[a, b, c] \mid \exists a_1, c_1, c_2 (R(a, b, c_1) \wedge R(a_1, 1, c) \wedge R(a_1, b, c_2))\}$$

##### 1.b)

$T_1$			
$a$	$b$	$c_1$	<b>R</b>
$a_1$	$1$	$c$	<b>R</b>
$a_1$	$b$	$c_2$	<b>R</b>

##### 2)

$$h_1 : T_2 \rightarrow T_1 : a \rightarrow a, b \rightarrow b, a_5 \rightarrow a_1, b_5 \rightarrow b_1, c_4 \rightarrow c_1 \\ \Rightarrow T_2 \subseteq T_1 \quad (1)$$

$$h_2 : T_1 \rightarrow T_2 : a \rightarrow a, b \rightarrow b, b_1 \rightarrow b, c_1 \rightarrow c_4, a_1 \rightarrow a, b_2 \rightarrow b_5, c_2 \rightarrow c_4, b_3 \rightarrow b_5, c_3 \rightarrow c_4 \\ \Rightarrow T_1 \subseteq T_2 \quad (2)$$

$$(1) \& (2) \Rightarrow T_1 \equiv T_2$$

## Exercise 3.2

Given 16 buffer pages (B) and :

- Album has a size of 10.000 pages (M), 40 bytes record size ( $s_1$ ) and 100 tuples/page ( $p_A$ )
- Track has a size of 200.000 pages (N), 30 bytes record size ( $s_2$ ) and 80 tuples/page ( $p_T$ )

1)

Since the simple nested loop join is a double iteration over both relations the I/O requirements can be calculated as follows:

$$M + p_A \cdot M \cdot N = 10.000 + 100 \cdot 10.000 \cdot 200.000 = 200.000.010.000 \text{ I/Os}$$

2)

Since the block nested loop join uses 1 input and 1 output buffer the number of I/Os can be calculated with the following formula:

$$M + \lceil \frac{M}{B-2} \rceil \cdot N = 10.000 + \lceil \frac{10.000}{16-2} \rceil \cdot 200.000 = 143.060.000$$

3)

Similarities:

- double loop schema

Differences:

- usage of buffer pages
- number of *outer* elements in the memory
- hashing used in the block nested loop

Explanation:

Lets define the outer loop to be over E and the inner loop to be over T.

The *block nested loop join* loads up to  $B - 2$  pages of E into the memory. These blocks get stored in a hash table. Now every input of a page from T can be used via one of the two reserved pages. After hash probing the input, if there is any match in a block of Es, we can reduce the number of scanned Es per T and therefore reduce I/Os.

## Exercise 3.3

1.

**Emp:**  $4.000 \text{ byte} / 20 \text{ byte} = 200 \Rightarrow 200 \text{ records fit on one page}$   
 $20.000 \text{ records} / 200 \frac{\text{records}}{\text{page}} = 100 \Rightarrow \text{need 100 pages}$

**DeptProj:**  $4.000 \text{ byte} / 40 \text{ byte} = 100 \Rightarrow 100 \text{ records fit on one page}$   
 $5.000 \text{ records} / 100 \frac{\text{records}}{\text{page}} = 50 \Rightarrow \text{need } 50 \text{ pages}$

**Proj:**  $4.000 \text{ byte} / 2.000 \text{ byte} = 2 \Rightarrow 2 \text{ records fit on one page}$   
 $1.000 \text{ records} / 2 \frac{\text{records}}{\text{page}} = 500 \Rightarrow \text{need } 500 \text{ pages}$

## 2.

Worst case: have to access all pages  $\xRightarrow{1} 100 \text{ pages}$

Index access costs 3  $\Rightarrow$  for  $N > 97$  scan is cheaper

## 3.

(a) Cost of a *block nested loop join*:  $M + \lceil \frac{M}{B-2} \rceil \cdot N$

$\xRightarrow{1} M = 50, N = 100, B = 12$ , so  
 $50 + \lceil \frac{50}{12-2} \rceil \cdot 10 = 550 \text{ I/Os}$

(b) Cost of a *index nested loop join*:  $M + ((M \cdot p_T) \cdot \text{cost of finding matching E tuples})$

$\xRightarrow{1} M = 50, p_T = 100$ , cost of finding matching E tuples =  $3 + 1 = 4$   
 $50 + ((50 \cdot 100) \cdot 4) = 20.050 \text{ I/Os}$

(c) The best option would be the *Sort-Merge Join* with costs of:  $M \log M + N \log N + (M + N)$   
 Since they are already sorted we can neglect the sorting cost of  $M \log M + N \log N$  and have remaining costs of  $M + N$

$\xRightarrow{1} M = 50, N = 100$   
 $50 + 100 = 150 \text{ I/Os}$

(d) Identical to 3.3.3.c, since the B+ tree is just a special form of sorting. So analogously we receive the costs  $M + N$

$\xRightarrow{1} M = 50, N = 100$   
 $50 + 100 = 150 \text{ I/Os}$