

## 6. Exercise

Timo Bergerbusch 344408

Thomas N  veke 311392

Shu Xing 381176

16.01.2018

### Exercise 6.1

#### 6.1.1

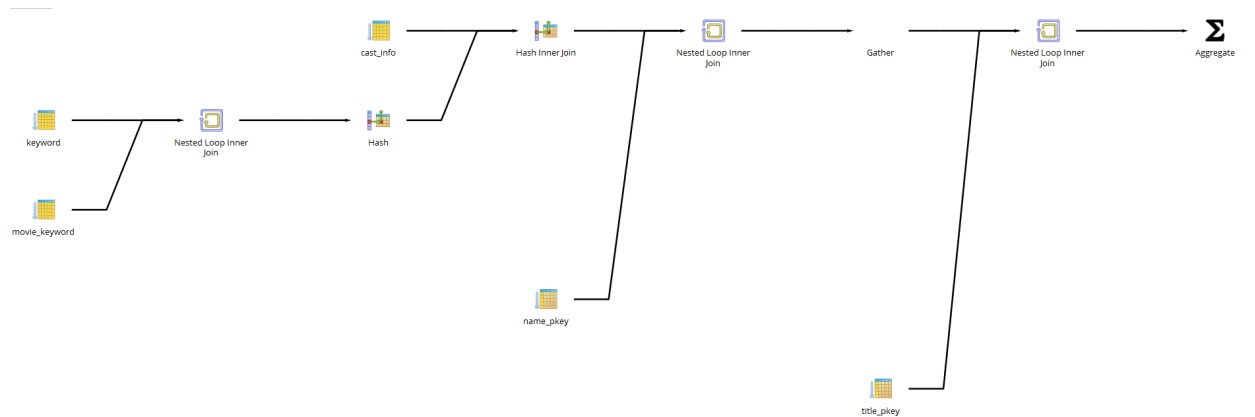


Figure 1: Execution Plan as calculated by PostgreSQL

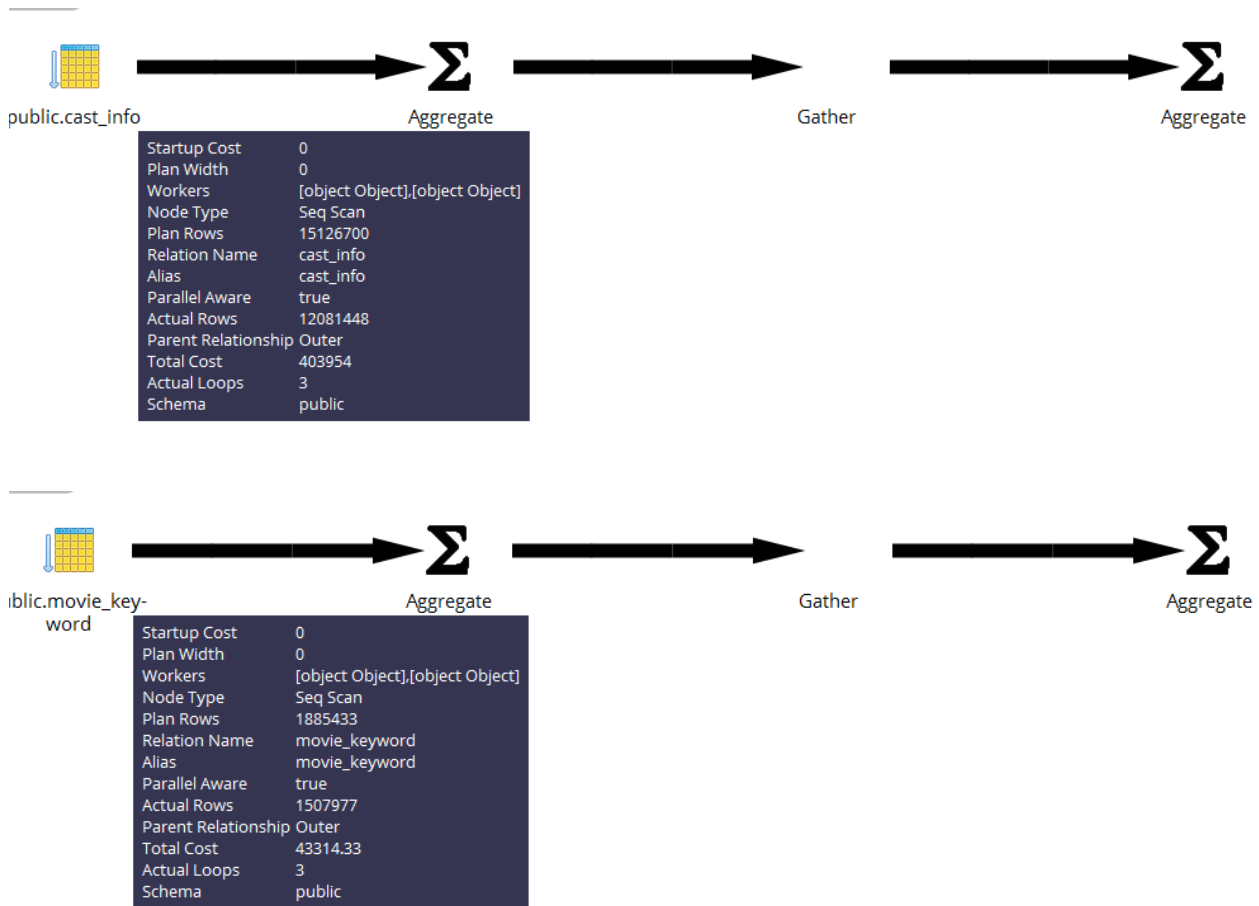
The most expensive operation will be scanning through *cast\_info* due to its sheer size.

#### 6.1.2

All cardinalities were determined with the query  
"EXPLAIN (FORMAT JSON, ANALYZE ON, VERBOSE ON, COSTS ON,  
BUFFERS OFF, TIMING OFF) SELECT COUNT(\*) FROM table\_name"

	estimated cardinality	true cardinality	q-error
* cast_info	15126700	12081448	1.252060184
name	1736764	1389164	1.250222436
* movie_keyword	1885433	1507977	1.250306205
keyword	135160	134170	1.007378699
title	1053407	842771	1.249932663

Figure 2: Cardinality and q-error



### 6.1.3

Since the proposed table differs from the query plan, we adjusted it to fit our query plan. For Screenshots see last page.

n.name LIKE '%Downey%Robert%'	1	0	0
k.keyword = 'marvel-cinematic-universe'	1	1	1
n.id = ci.person_id	1	2	2
k.id = mk.keyword_id	33	14	2.357142857
ci.movie_id = mk.movie_id	555	414	1.34057971
mk.movie_id = t.id	1	6	6

Figure 3: Cardinality and q-error

### 6.1.4

As we can read from the charts, the estimated and true cardinality often differ. This is due to the statistical sampling of several entries the optimizer uses to determine the estimated values. We could increase the accuracy by increasing the number of entries sampled. The closer we get to the number of entries, the closer the values will be to reality. However that makes the whole process of optimization very inefficient.

## 6.1.5

## Exercise 6.2 (Datalog)

### 6.2.1

a) `parent(X,Y) :- child(Y,X).`

b) `married(X,Y) :- child(Z,X), child(Z,Y).`

c) `sister(X,Y) :- child(X,Z), child(Y,Z), NOT male(X).`


Comment: assuming halvesisters are also considered sisters.

d) `halfbrother(X,Y) :- parent(Z,X), parent(Z,Y), parent(N,X), NOT parent(N,Y), male(X).`

Comment: assuming every person has only two parents.


### 6.2.2

- Datalog program 1:



No two predicates in a layer depend negatively on each other, so the program is stratified.

- Datalog program 2:



Since the predicate  $\tau$  depends negatively on itself program 2 is not stratified

Index Cond	(n.id = ci.person_id)	
Startup Cost	0.43	
Scan Direction	Forward	
Plan Width	19	
Rows Removed by Index Recheck	0	
Workers	[object Object],[object Object]	
Node Type	Index Scan	
Total Cost	0.46	
Plan Rows	1	
Relation Name	name	
Alias	n	
Parallel Aware	false	
Actual Rows	0	
Output	n.id,n.name,n.imdb_index,n.imdb_id,n.gender,n.name_pcode_cf,n.name_pcode_nf,surname_pcode,n.md5sum	
Parent Relationship	Inner	
Schema	public	
Filter	((n.name)::text -- '%Downey%Robert%':text)	
Actual Loops	1242	
Rows Removed by Filter	1	
Index Name	name_pkey	

Filter	((k.keyword)::text = 'marvel-cinematic-universe':text)	
Startup Cost	0	
Plan Width	20	
Workers	[object Object],[object Object]	
Node Type	Seq Scan	
Plan Rows	1	
Relation Name	keyword	
Alias	k	
Parallel Aware	false	
Actual Rows	1	
Output	k.id,k.keyword,k.phonetic_code	
Parent Relationship	Outer	
Total Cost	2645.5	
Actual Loops	3	
Rows Removed by Filter	134169	
Schema	public	